

Project Overview

STEPS

- Data Collection [kaggle](#)
- Feature Engineering
- Feature Scaling
- Model Creation
- Evolution of Model

Attribute Information:

Input variables (based on physicochemical tests):

1.fixed acidity 2.volatile acidity 3.citric acid 4.residual sugar 5.chlorides 6.free sulfur dioxide
7.total sulfur dioxide 8.density 9.pH 10.sulphates 11.alcohol

1. quality (score between 0 and 10) Output variable (based on sensory data):

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')

df = pd.read_csv('/Data/winequalityN.csv')
```

Exploratory Data Analysis

```
df.head()
```

	type	fixed acidity	volatile acidity	citric acid	residual sugar
0	white	7.0	0.27	0.36	20.7
1	white	6.3	0.30	0.34	1.6
2	white	8.1	0.28	0.40	6.9
3	white	7.2	0.23	0.32	8.5
4	white	7.2	0.23	0.32	8.5

	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH
0	0.045	45.0	170.0	1.0010	3.00

1	0.049	14.0	132.0	0.9940	3.30
2	0.050	30.0	97.0	0.9951	3.26
3	0.058	47.0	186.0	0.9956	3.19
4	0.058	47.0	186.0	0.9956	3.19

	sulphates	alcohol	quality
0	0.45	8.8	6
1	0.49	9.5	6
2	0.44	10.1	6
3	0.40	9.9	6
4	0.40	9.9	6

```
from sklearn.preprocessing import LabelEncoder
label = LabelEncoder()
df['type']=label.fit_transform(df['type']) #white=1,red=0
```

```
## info about dataset
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6497 entries, 0 to 6496
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   type                                  6497 non-null   int32
1   fixed acidity                         6487 non-null   float64
2   volatile acidity                      6489 non-null   float64
3   citric acid                           6494 non-null   float64
4   residual sugar                        6495 non-null   float64
5   chlorides                             6495 non-null   float64
6   free sulfur dioxide                   6497 non-null   float64
7   total sulfur dioxide                  6497 non-null   float64
8   density                              6497 non-null   float64
9   pH                                    6488 non-null   float64
10  sulphates                             6493 non-null   float64
11  alcohol                               6497 non-null   float64
12  quality                               6497 non-null   int64
dtypes: float64(11), int32(1), int64(1)
memory usage: 634.6 KB
```

```
## size of the dataframe
df.shape
```

```
(6497, 13)
```

```
## duplicated rows
df.duplicated().any()
```

True

dropping duplicate

```
df.drop_duplicates(inplace=True)
```

df.shape ## shape of df after removing the dataset

(5329, 13)

Missing values


```
df.isnull().sum()
```

type	0
fixed acidity	10
volatile acidity	8
citric acid	3
residual sugar	2
chlorides	2
free sulfur dioxide	0
total sulfur dioxide	0
density	0
pH	9
sulphates	4
alcohol	0
quality	0

dtype: int64

```
sns.heatmap(df.isnull())
```

<Axes: >



- columns are fixed acidity,volatile acidity,citric acid,residual sugar,chlorides,pH,sulphates are have the some of null values init

Handling Missing Values

```
## missing columns
columns_having_missing = {}
for column in df.columns:
    if df[column].isnull().sum() > 0:
        columns_having_missing[column] = df[column].isnull().sum()
print(columns_having_missing)

{'fixed acidity': 10, 'volatile acidity': 8, 'citric acid': 3,
 'residual sugar': 2, 'chlorides': 2, 'pH': 9, 'sulphates': 4}
```

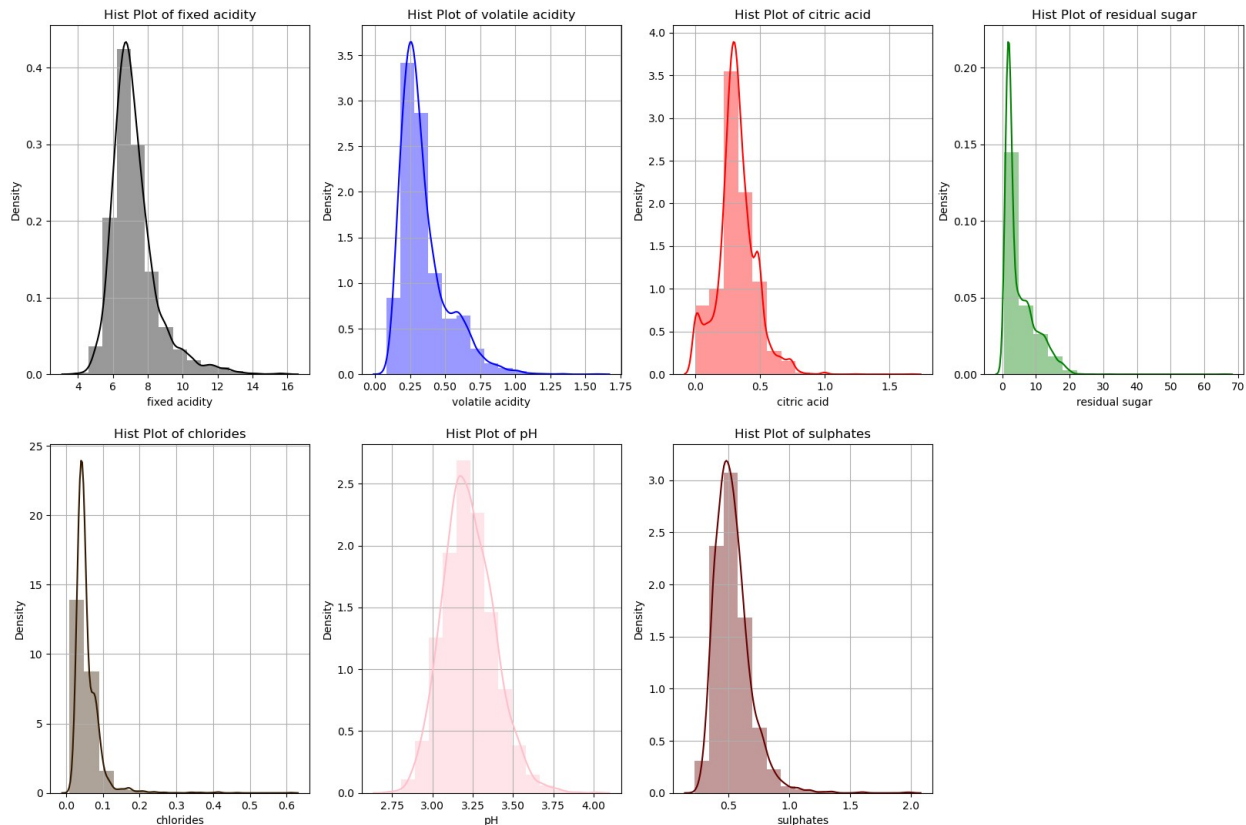
histplots before filling the miss values in columns

```
## histplots before filling the miss values in columns
pt.figure(figsize=(20, 20))
```

```

color = ['k','b','r','g','#331900','pink','#660000']
for i, feature in enumerate(list(columns_having_missing.keys())):
    pt.subplot(3,4 ,i + 1)
    sns.distplot(df[feature],color=color[i],bins=15)
    pt.grid(axis='both')
    pt.title(f'Hist Plot of {feature}')

```



```

##Filling the missing values with its column's mean
for feature in columns_having_missing.keys():
    df[feature].fillna(df[feature].mean(),inplace=True)

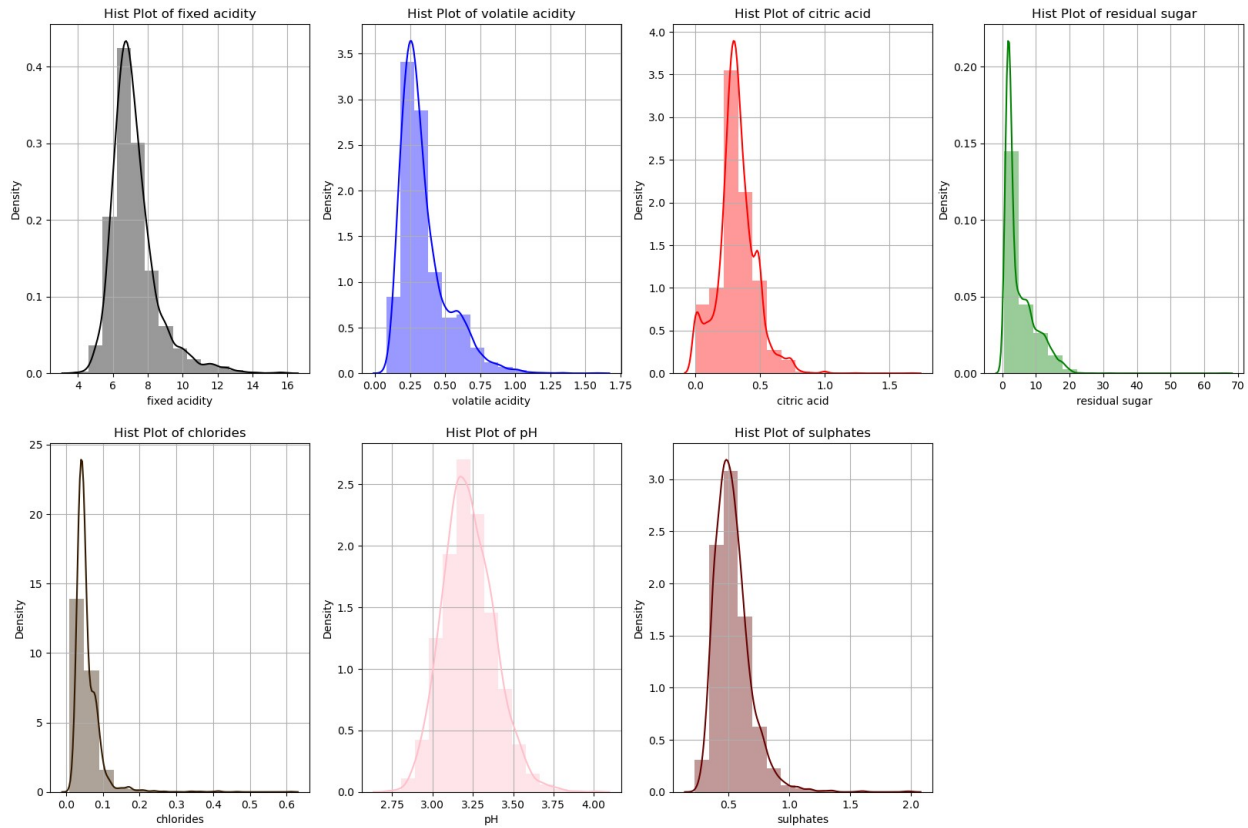
```

histplots after filling the miss values in columns

```

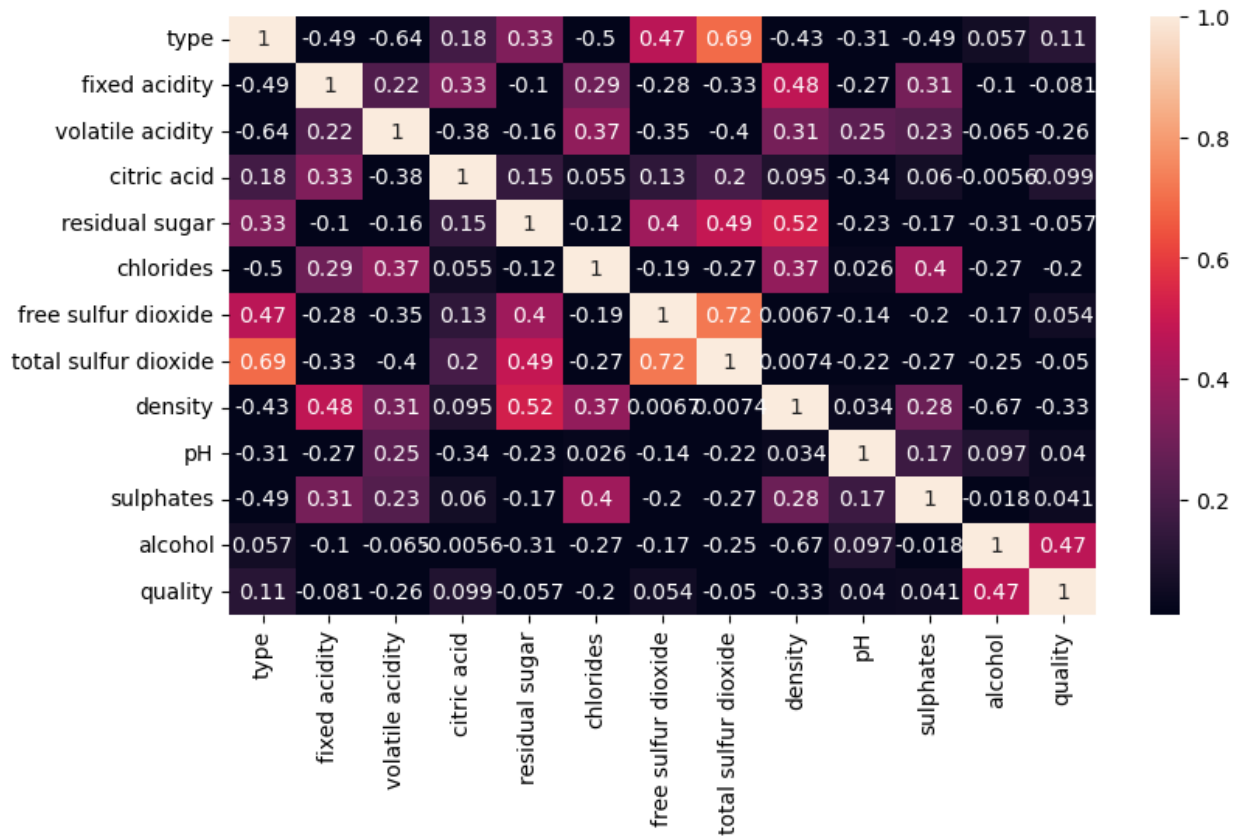
## seeing the distribution after filling the histplots
pt.figure(figsize=(20, 20))
color = ['k','b','r','g','#331900','pink','#660000']
for i, feature in enumerate(list(columns_having_missing.keys())):
    pt.subplot(3,4 ,i + 1)
    sns.distplot(df[feature],color=color[i],bins=15)
    pt.grid(axis='both')
    pt.title(f'Hist Plot of {feature}')

```

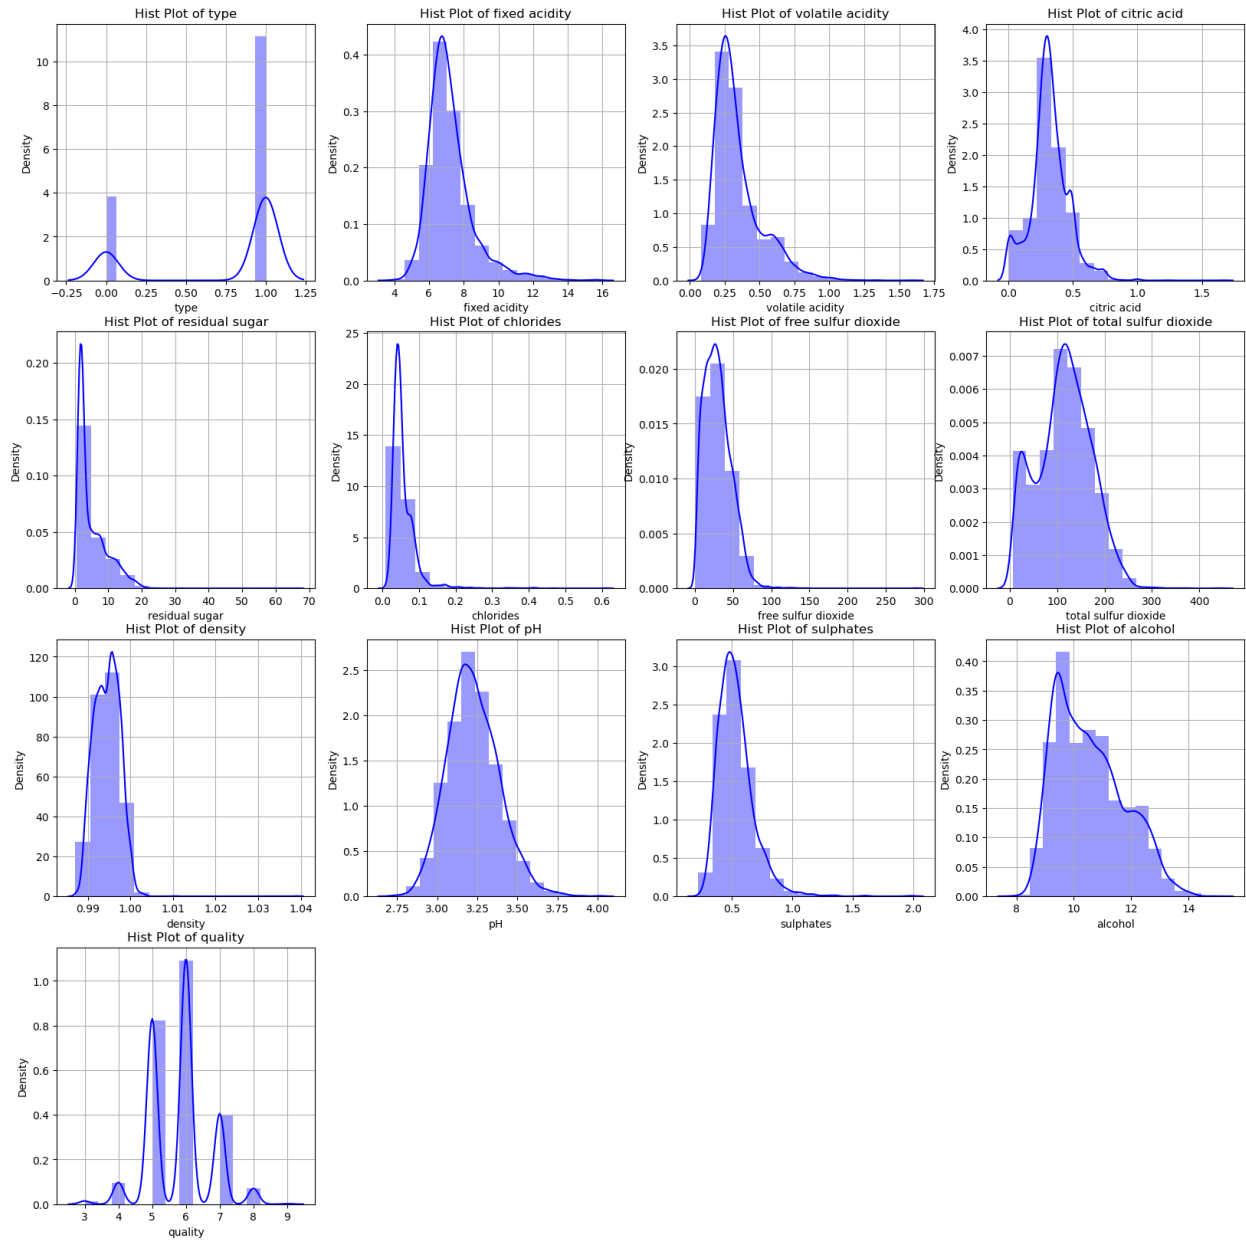


```
## correlation matrix
plt.figure(figsize=(9,5))
sns.heatmap(df.corr(),annot=True,vmin=0.01)
```

<Axes: >



```
## Univaraity analysis
pt.figure(figsize=(20, 20))
for i, feature in enumerate(df.columns):
    pt.subplot(4,4 ,i + 1)
    sns.distplot(df[feature],color='blue',bins=15)
    pt.grid(axis='both')
    pt.title(f'Hist Plot of {feature}')
```



```
df.head()
```

	type	fixed acidity	volatile acidity	citric acid	residual sugar
0	1	7.0	0.27	0.36	20.7
1	1	6.3	0.30	0.34	1.6
2	1	8.1	0.28	0.40	6.9
3	1	7.2	0.23	0.32	8.5
6	1	6.2	0.32	0.16	7.0

	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH
\					
0	0.045	45.0	170.0	1.0010	3.00
1	0.049	14.0	132.0	0.9940	3.30
2	0.050	30.0	97.0	0.9951	3.26
3	0.058	47.0	186.0	0.9956	3.19
6	0.045	30.0	136.0	0.9949	3.18

	sulphates	alcohol	quality
0	0.45	8.8	6
1	0.49	9.5	6
2	0.44	10.1	6
3	0.40	9.9	6
6	0.47	9.6	6

Bivariate Analysis

```

pt.figure(figsize=(20,20))
pt.title('Bivariat Analysis')
pt.subplot(3,3,1)
sns.scatterplot(df,x='fixed acidity',y='volatile
acidity',color='g',hue='quality',palette=['black','red','green','pink',
'yellow','brown','blue'])
pt.title('fixed acidity vs volatile acidity')

pt.subplot(3,3,2)
sns.scatterplot(df,x='free sulfur dioxide',y='total sulfur
dioxide',hue='quality',palette='bright')
pt.title('free sulfur dioxide vs total sulfur dioxide')

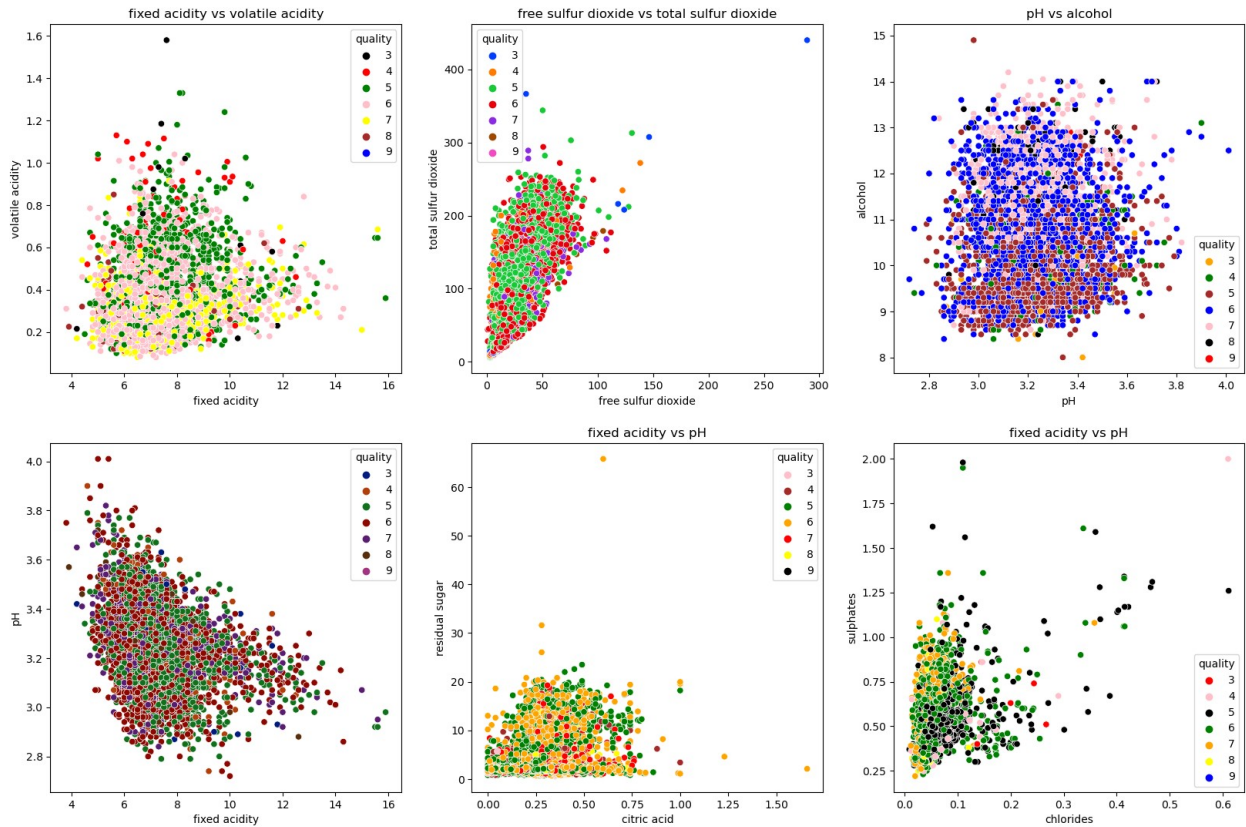
pt.subplot(3,3,3)
sns.scatterplot(df,x='pH',y='alcohol',hue='quality',palette=['orange',
'green','brown','blue','pink','black','red'])
pt.title('pH vs alcohol')

pt.subplot(3,3,4)
sns.scatterplot(df,x='fixed
acidity',y='pH',hue='quality',palette='dark')

pt.subplot(3,3,5)
sns.scatterplot(df,x='citric acid',y='residual
sugar',hue='quality',palette=['pink','brown','green','orange','red','y
ellow','black'])
pt.title('fixed acidity vs pH')

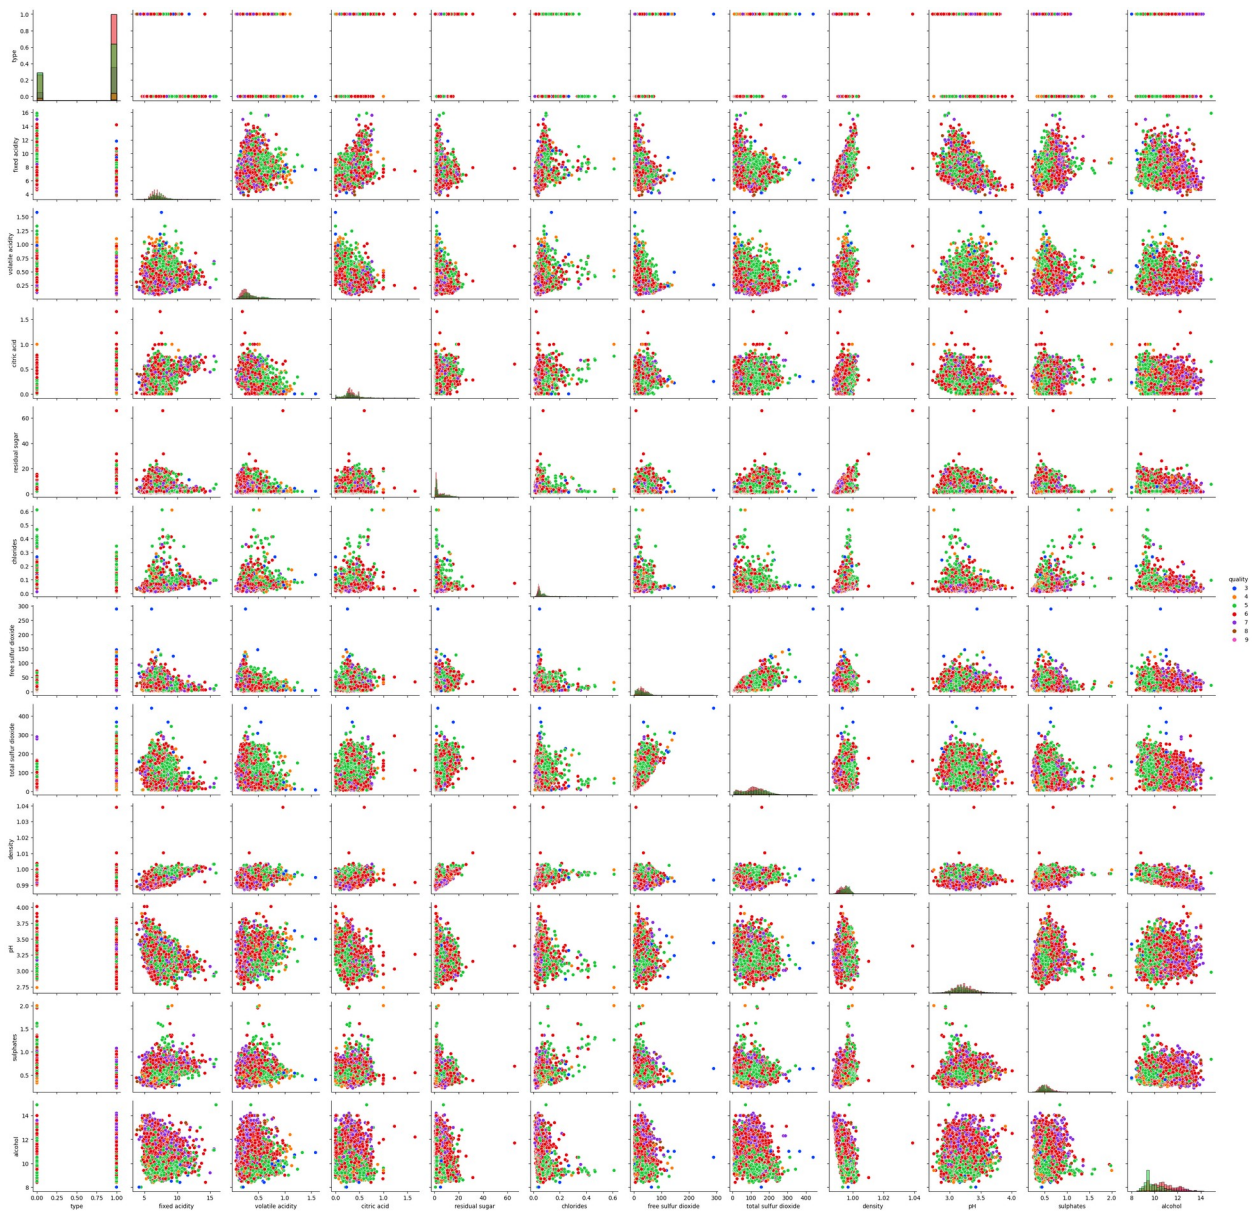
```

```
pt.subplot(3,3,6)
sns.scatterplot(df,x='chlorides',y='sulphates',hue='quality',palette=[
'red','pink','black','green','orange','yellow','blue'])
pt.title('fixed acidity vs pH')
pt.show()
```




```
g = sns.PairGrid(df, hue="quality",palette='bright')
g.map_diag(sns.histplot)
g.map_offdiag(sns.scatterplot)
g.add_legend()
```

<seaborn.axisgrid.PairGrid at 0x19c0fabf290>



Outliers

```
pt.figure(figsize=(20, 20))
colors =
['pink', 'brown', 'green', 'orange', 'red', 'yellow', 'black', 'green', 'brown',
', 'blue', 'pink', 'black', 'red', 'yellow']
for i, feature in enumerate(df.columns):
    pt.subplot(4,4 ,i + 1)
    sns.boxplot(df[feature],color=colors[i+1])
    pt.grid(axis='both')
    pt.title(f'Box Plot of {feature}')
```



Mapping the target output classes as Low,Medium,High Quality

Target column have 3,4,5,6,7,8,9 are output classes Generaizing output classes

- (3,4) --> LOW Quality
- (5,6,7) --> MEDIUM Quality
- (8,9) --> HIGH Quality

Mapping to output classes

```
df['quality'] =
df['quality'].map({3:'Low',4:'Low',5:'Medium',6:'Medium',7:'Medium',8:
'High',9:'High'})
```

```
df['quality'] = df['quality'].map({'Low':0,'Medium':1,"High":2})
```

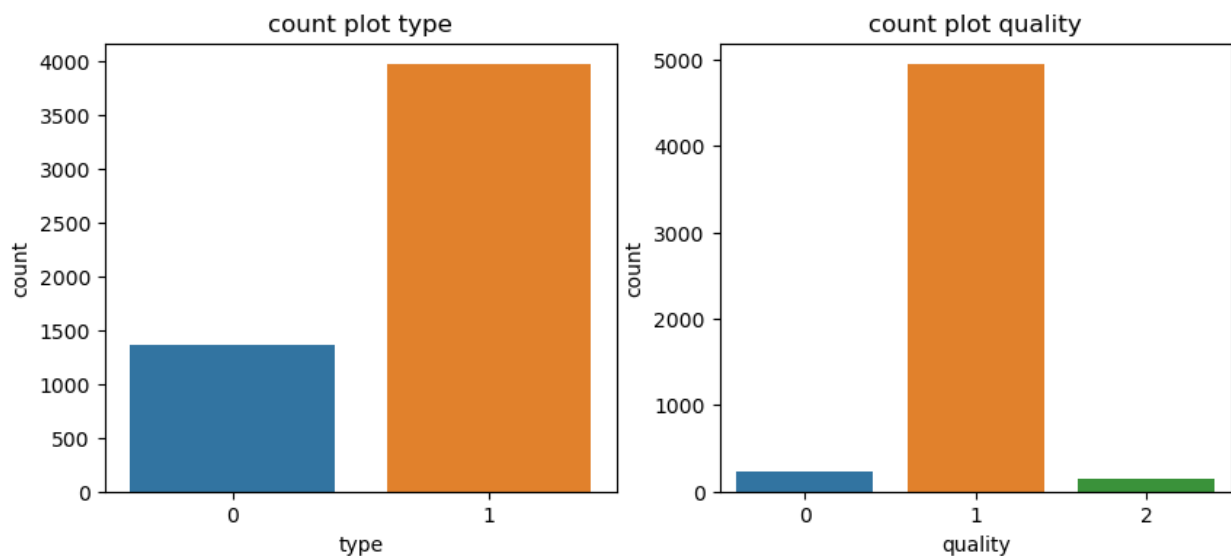
```

df.quality.unique()

array([1, 2, 0], dtype=int64)

## categorical feature count plot
pt.figure(figsize=(10,4))
for i,feature in enumerate(df[['type','quality']].columns):
    pt.subplot(1,2,i+1)
    sns.countplot(df,x=feature)
    pt.title('count plot '+feature)

```



```

##checking the data is whether is imbalanced dataset
print('Low',len(df[df['quality']==0]))
print('Medium',len(df[df['quality']==1]))
print('High',len(df[df['quality']==2]))

Low 236
Medium 4939
High 154

X = df.drop('quality',axis=1)
y = df.quality

print(X.shape)
print(y.shape)

(5329, 12)
(5329,)

```

Feature Importance

```

#### Feature Importance
from sklearn.ensemble import ExtraTreesClassifier

```

```

feature_imp= ExtraTreesClassifier()
feature_imp.fit(X,y)
scores = feature_imp.feature_importances_
pd.DataFrame({'Feature': X.columns, 'Feature_Importance':
(scores*100)}).sort_values(by = 'Feature_Importance', ascending =
True)

```

	Feature	Feature_Importance
0	type	1.024119
8	density	8.073902
5	chlorides	8.266900
3	citric acid	8.410160
1	fixed acidity	8.463813
9	pH	8.499229
4	residual sugar	8.573696
10	sulphates	8.792134
7	total sulfur dioxide	8.861663
11	alcohol	9.442263
2	volatile acidity	10.767609
6	free sulfur dioxide	10.824512

Converting the Imbalanced dataset into Balanced dataset

```

# transform the dataset in to balanced format
from imblearn.over_sampling import SMOTE
oversample = SMOTE(k_neighbors=4)
X, y = oversample.fit_resample(X, y)

print(X.shape)
print(y.shape)

(14817, 12)
(14817,)


y.value_counts()

1    4939
2    4939
0    4939
Name: quality, dtype: int64

### The target classes in balanced format
y_df = pd.DataFrame(np.array(y), columns=['target'])
sns.countplot(y_df, x='target')
plt.ylim([1000, 5000])

(1000.0, 5000.0)

```



```
df.head()
```

	type	fixed acidity	volatile acidity	citric acid	residual sugar
0	1	7.0	0.27	0.36	20.7
1	1	6.3	0.30	0.34	1.6
2	1	8.1	0.28	0.40	6.9
3	1	7.2	0.23	0.32	8.5
6	1	6.2	0.32	0.16	7.0
	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH
0	0.045	45.0	170.0	1.0010	3.00
1	0.049	14.0	132.0	0.9940	3.30
2	0.050	30.0	97.0	0.9951	3.26
3	0.058	47.0	186.0	0.9956	3.19

6	0.045	30.0	136.0	0.9949	3.18
---	-------	------	-------	--------	------

	sulphates	alcohol	quality
0	0.45	8.8	1
1	0.49	9.5	1
2	0.44	10.1	1
3	0.40	9.9	1
6	0.47	9.6	1

Train Test Split

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test =
train_test_split(X,y,test_size=0.2,random_state=42)
```

Feature Scaling

Model Training

```
##### Model selection
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier,AdaBoostClassifier
from xgboost import XGBClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import cross_val_score
```

```
models = {'Logistic Regression':LogisticRegression(),
          'Support Vector Machine':SVC(),
          'DecsionTree':DecisionTreeClassifier(),
          'RadomForestClassier':RandomForestClassifier(),
          'AdaBosster':AdaBoostClassifier(),
          'XGBboost':XGBClassifier()}
```

```
for i in range(len(models)):
    model =list(models.values())[i]
    model.fit(X_train,y_train)
    print(list(models.keys())[i]+' score:
',model.score(X_test,y_test))
    cros_score = cross_val_score(model,X_train,y_train,cv=5)
    print(list(models.keys())[i]+' Cross_Val :',list(cros_score))
    print('mean : ',np.mean(cros_score))
    print('___'*40)
```

```
Logistic Regression score: 0.6835357624831309
Logistic Regression Cross_Val : [0.6596372838464782,
0.6625896246309574, 0.6748207507380852, 0.6620253164556962,
0.6662447257383967]
mean : 0.6650635402819227
```



```
Support Vector Machine score: 0.8488529014844804
Support Vector Machine Cross_Val : [0.8304512863770561,
0.8253901307465205, 0.8253901307465205, 0.8227848101265823,
0.8265822784810126]
mean : 0.8261197272955384
```

```
DecsionTree score: 0.9024966261808367
DecsionTree Cross_Val : [0.8857022353437368, 0.8975115984816533,
0.8810628426824125, 0.8662447257383966, 0.870042194092827]
mean : 0.8801127192678052
```

```
RadomForestClassier score: 0.9622132253711201
RadomForestClassier Cross_Val : [0.9603542808941375,
0.9502319696330662, 0.956980177140447, 0.9531645569620253,
0.9514767932489452]
mean : 0.9544415555757242
```

```
AdaBosster score: 0.6835357624831309
AdaBosster Cross_Val : [0.6992830029523408, 0.6840995360607338,
0.6811471952762548, 0.6734177215189874, 0.6649789029535865]
mean : 0.6805852717523807
```

```
XGBboost score: 0.9723346828609987
XGBboost Cross_Val : [0.972163644032054, 0.9734289329396879,
0.9730071699704765, 0.9729957805907173, 0.970042194092827]
mean : 0.9723275443251527
```

Here RandomForest,XGBbooster are giving more score compare to other models so I am taking XGBbooster as my final model and doing hyperparameter tuning on it

HYPER PARAMETER TUNNING

XGBooster Tunning

```
Xgb = {
    'n_estimators': [50, 100, 200],
    'learning_rate': [0.01, 0.1, 0.5],
    'max_depth': [3, 5, 7],
    'min_child_weight': [1, 3, 5],
}

from sklearn.model_selection import GridSearchCV,RandomizedSearchCV
dic = {}
```

```

grid = GridSearchCV(XGBClassifier(),param_grid=Xgb,cv=5)
grid.fit(X_train,y_train)
dic['xgb'] = grid.best_params_

xgb = XGBClassifier(learning_rate=0.5,max_depth= 7,min_child_weight=
1,n_estimators=200)
xgb.fit(X_train,y_train)
xgb.score(X_test,y_test)

0.9794197031039136

## freecodecamp.org is the providing cources with certificates

###Test Data
y_pred = xgb.predict(X_test)

```

Model Evolution Metrics

```

### Perfomence metrics
from sklearn.metrics import
f1_score,classification_report,accuracy_score,recall_score,precision_s
core,confusion_matrix,auc
print('accuracy: ',accuracy_score(y_pred,y_test))
print('recall: ',recall_score(y_pred,y_test,average=None))
print('precision: ',precision_score(y_pred,y_test,average=None))
print('classification report: ',classification_report(y_pred,y_test))

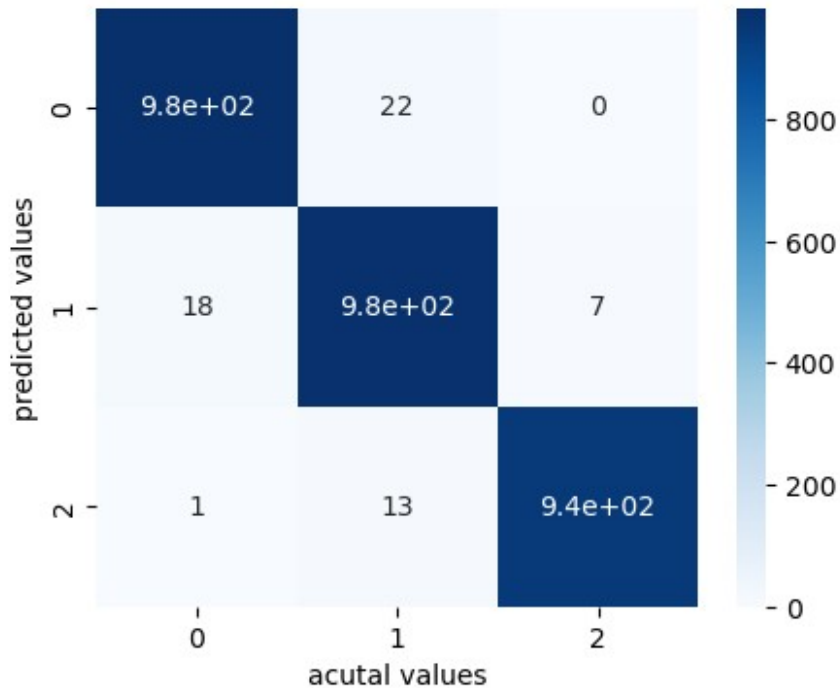
accuracy:  0.9794197031039136
recall:   [0.97810945 0.97507478 0.98535565]
precision: [0.98103792 0.96544916 0.99262381]
classification report:
precision    recall  f1-score   support

      0       0.98       0.98       0.98        1005
      1       0.97       0.98       0.97        1003
      2       0.99       0.99       0.99         956

   accuracy          0.98          0.98          0.98          2964
  macro avg          0.98          0.98          0.98          2964
weighted avg          0.98          0.98          0.98          2964

##Heatmap
pt.figure(figsize=(5,4))
sns.heatmap(confusion_matrix(y_pred,y_test),annot=True,cmap='Blues')
pt.xlabel('acutal values')
pt.ylabel('predicted values')
pt.show()

```



```
accuracy_dataframe3 = pd.DataFrame({"y_test": y_test, "y_pred":
y_pred})
```

```
print( 'Acutally points are incorrectly classified',
sum(accuracy_dataframe3['y_test']-
accuracy_dataframe3['y_pred']), 'points')
```

Acutally points are incorrectly classified -4 points

pickle file

```
###pickle file
import pickle
with open('model.pkl', 'wb') as files:
    pickle.dump(model, files)

with open('model.pkl' , 'rb') as f:
    model = pickle.load(f)

model.predict([[1,7.0,0.270,0.36,20.7,0.045,45.0,170.0,1.00100,3.00,0.
450000,8.8]])

array([1], dtype=int64)
```